Reversium by 3Sci Labs

Start here

What Reversium helps you do

Reversium reads a software package and organizes supported source and binary clues into findings, diagrams, evidence and explicit gaps. Use it when receiving a handover, inheriting unfamiliar code, investigating a feature or preparing maintenance. Ordinary browser analysis does not install, run or certify software; the separate experimental local sandbox tests one reviewed build/start profile.

Choose one useful outcome

For an owner: identify what was delivered and agree on one feature to check. For a developer: find likely entry points, dependencies and source locations. For a maintainer: map possible affected components before changing them. A useful outcome is a short, evidence-backed checklist rather than a claim that the entire platform is understood.

A first session in five steps

1. Open the browser version and choose Start exploring. 2. Choose Understand the software, Prepare a developer handover, or Plan maintenance or modernization. 3. Continue to ZIP, select your source-code package and choose Analyze ZIP locally. 4. Read Review findings, then inspect one Architecture workflow and its Evidence. 5. Read Coverage and Save & next steps; save a report and assign the remaining checks.

Finding help

Help is available on Start and in the report tabs. Search by task or keyword, such as handover, evidence or sandbox. Matching chapters open automatically; clear the search to show every chapter again. Download user guide (PDF) keeps an offline copy. The guide shows its build number and is rebuilt with each release.

Example questions

Where does the save action lead? Which services are referenced? Which files should a new developer read first? What configuration, permissions and sample data still need confirmation? Reversium provides investigation leads; isolated installation and runtime tests answer whether those leads work.

Online and local editions

Online browser edition

Open the browser edition from the Reversium product page. No installation, account or payment is needed for analysis. The app downloads its self-hosted Python runtime, then reads the selected ZIP in a browser worker on your device. Use a current desktop browser. The first visit downloads about 14 MB of app/runtime assets and can take longer.

Local application

Use the local application package supplied by 3Sci Labs. Install Python 3.10 or newer, open a terminal in the project and run python -m reversium. Open the local address printed in the terminal; it can use a free port. On Windows, Start-Reversium.ps1 is an alternative launcher. Stop the server with Ctrl+C. The base analyzer has no third-party Python dependency; optional AI has separate requirements.

Main differences

Both editions use the same analysis engine, findings and core diagrams. Online: 32 MiB compressed ZIP limit, five-minute worker stop, browser-memory constraints, no embedded cloud account connection. Local web application: 64 MiB upload limit, local Python process and existing optional ChatGPT account connection. Local command line: report exports and baseline comparison without a web page; archive safety limits still apply.

AI differences

Online AI is a manual handoff: select a provider and evidence, review/edit the exact request, download it, then decide whether to paste and submit it on that provider website. It is not an account connection or embedded chat. Local AI offers an account connection inside Reversium; install requirements-ai.txt if needed, follow its connection wizard and review sharing before each request. Eligibility, network access and provider limits apply. A successful sign-in does not guarantee a successful model response.

Which should I use?

Start online for a quick inspection without installing anything. Use local for larger supported packages, repeatable command-line work or the existing optional account integration. Use neither as proof that a platform is production-ready. Online remains free to analyze; optional AI services follow the terms and limits of the account you choose.

Selecting a package and handling files

Choose the best input

Prefer the project source-code package with manifests and migrations. An installer or binary-only archive provides fewer clues. Do not include unrelated backups, generated dependencies or large media if they are unnecessary. Keep the original delivery intact and analyze a separate copy when preparing a smaller archive.

What stays on the device

The online app does not upload the ZIP. It holds the file and report in page/worker memory, not app-managed localStorage or IndexedDB. Static web hosting receives normal app and asset requests, including connection information such as IP address; those requests do not include ZIP contents. No input source is executed. No AI evidence is submitted automatically.

Clearing and saving

Cancel and clear stops the worker and returns to Start. Clear and start again removes the report and selected-file references. Closing the page also releases app references; this is not a secure-memory-erasure guarantee. Browser caches may retain application assets. Downloaded reports and requests remain wherever you save them and can contain private filenames or code-derived details. Share them deliberately.

Limits and rejected archives

Online input is limited to a non-empty ZIP no larger than 32 MiB; the local web application accepts up to 64 MiB. Analyzer bounds include 10,000 files, 256 MiB expanded total and 4 MiB content inspection per file, plus compression, path and graph bounds. Encrypted, damaged, unsafe or oversized archives may be rejected or only partly inspected. Nested containers receive bounded directory inspection, not recursive analysis. Read the actual Coverage output.

Reading the main report

Header and inventory

The logo and build number identify the report generator. Main tabs stay near the top. Technical inventory numbers remain expanded: files, bytes, binary candidates and findings. These are counts of inspected material, not a completeness, quality or security score.

About and Review findings

About explains Reversium itself. Review findings describes supported clues about the uploaded software in plain language. Open Why we say this to inspect the evidence behind a statement. Detailed findings for a developer expands more technical explanations. Different starting goals alter guidance, not the underlying source evidence.

Evidence

Use the search box to filter findings by file, symbol, value, kind or resolution. A source link from a feature opens Evidence and reveals its row. Locations are file and line references, or byte offsets for binary clues; they are not an in-browser source editor. Traced means a supported static reference was found. Inferred means a candidate relationship needs confirmation. Neither means the program ran successfully.

What to do with unknowns

Record the question, the relevant finding, the person who can resolve it, and a concrete verification step. For example: confirm the API route and database table by running one isolated user journey with test data and appropriate permissions. Missing references can reflect unsupported syntax or missing files, not missing functionality.

Architecture and relationship diagrams

Subsection buttons

Open Architecture and choose Workflow, Graphical feature map, Feature catalogue, Platform profile, Supporting services, or Imports & dependencies. The selected subsection replaces the other views, so you do not need to scroll through the whole architecture report. Available sections depend on report content.

Workflow

Choose a screen and data branch. Read the ordered path from screen through callbacks and an API to a database operation. Each step links to its evidence. The view follows a bounded candidate path; it can show a logging side effect rather than the main purpose of a request. It does not establish runtime conditions, arguments, permissions or successful responses. Read What remains unknown.

Graphical feature map

Select a feature to focus the map. Click a node or connection to inspect participants, evidence and unknowns. Solid lines represent traced source references; dashed lines represent inferred connections. Shared names do not prove a shared database. The display is capped at 120 nodes and 240 edges per view, with omissions disclosed. Download map SVG saves the drawn map; JSON retains the full generated graph, which is itself bounded.

Feature catalogue and file relationships

The catalogue lists supported feature candidates, operations and evidence. Imports & dependencies draws file-to-local-module candidates or referenced packages; it shows the first 100 unique edges and is not a complete runtime call graph. Full available findings remain in Evidence. Platform profile lists candidate technologies and runtimes. Supporting services lists references that may be optional, development-only or environment-dependent.

When a map looks incomplete

Filter to one feature and inspect its evidence. Check Coverage for unresolved targets and unsupported syntax. Dynamic code, aliases, middleware, component instances, conditions and runtime configuration can break a static path. Do not infer that an absent edge proves there is no relationship. Large diagrams are an investigation aid, not a complete specification.

Coverage, binary clues and readiness

Coverage subsection buttons

Overview describes what inspection can establish. Binary inspection shows supported metadata readers and target/runtime clues. File inventory lists path, type, byte size and inspection method. Setup & readiness lists questions and suggested checks. Limitations records warnings and analysis boundaries specific to the package.

Supported source analysis

Python uses AST parsing for supported imports, symbols and selected route/data traces. JavaScript/TypeScript and React use bounded token/reference analysis for selected routes, callbacks, prop bindings and computed API targets. Supabase table, RPC, Auth and storage references and selected SQL migration links produce candidates. Other languages and unsupported patterns may receive inventory-only coverage. Inspect the report rather than assuming a language-wide guarantee.

Binary analysis

Binary inspection is bounded static metadata and reference analysis, not full decompilation or behavioral reconstruction. A binary reader finding a format, target or string does not prove that the executable is safe, valid or runnable. Nested APK/JAR/Office/archive containers are not automatically reconstructed into a complete application.

Readiness is a separate test

Reversium does not install dependencies, start services, connect to production, validate database permissions or test a live user journey. To verify a handover, use an isolated environment, disposable accounts and sample data. Record required runtimes, environment-variable names, migrations, storage, integrations, permissions and successful/failed steps. Never assume a referenced cloud service is configured or paid for.

Optional AI, downloads and comparison

Online AI handoff

Open AI / choose provider after analysis. Choose ChatGPT, Claude, Gemini or Another AI. Enter a question, search findings and tick up to 20 findings. Prepare request for review builds an editable request with IDs, values, locations, confidence and resolution. Review exactly what it includes; remove secrets or private details. Download reviewed request saves text locally. Opening the provider website sends no request text; you must paste and submit it yourself. No AI answer is displayed inside the browser app.

Checking an AI answer

Ask for finding IDs and locations, facts separated from inferences, and explicit unknowns. Reopen cited evidence and check whether the answer follows from the selected subset. Reject invented runtime claims, missing citations or claims of complete coverage. Source-derived text can contain misleading instructions; treat it as untrusted evidence. Provider terms and account limits apply when you submit content there.

Save readable report: HTML

Use Save & next steps, then Save readable report (HTML). Open the saved file later without running Reversium, or share it with an authorized colleague. It retains generated findings, diagrams and limits; it is not the original ZIP or an application backup. Account credentials and the assistant conversation are excluded. Keep a separate copy of an AI request if you need its exact wording.

Save technical data: JSON and map SVG

JSON is a structured record for scripts, inspection and comparing deliveries. It includes the generated analysis and full bounded graph, not original source files. SVG is useful for documents or discussions about the drawn map. Neither file rebuilds or runs the application. Save both HTML and JSON as a baseline when accepting a delivery.

Browser comparison

Open Compare reports and select an earlier report JSON, up to 32 MiB. File additions/removals, shared files with different hashes, finding additions/removals and shared files without comparable hashes are counted. Line movement alone does not change finding identity. Lists show up to 200 entries each; Save comparison notes preserves the displayed results and omissions. Feature-map changes also compare recorded nodes, features and relationships; see the feature-map comparison chapter. Reports from different analyzer versions can differ without a code change. Comparisons do not overwrite checklist notes or enter AI packets. No files are uploaded.

Local command-line comparison

Run python -m reversium project.zip --output reports/project to write report.html, report.json and feature-map.svg. Add --compare previous-report.json to compare supported file/finding and feature-map identities. Outputs overwrite the same filenames in the selected directory. Structural differences and content hashes do not establish a behavioral change; old baselines may lack map fields.

Handover and change-impact review

Handover tab

Open Handover after analysis. Expand Setup clues to inspect referenced environment variable names, names from .env.example/.env.sample/.env.template, package.json script names and declared runtime constraints. Follow their evidence links. The reader does not retain environment values or script commands; names may occur in comments or optional code and do not prove a requirement. Actual .env files are not read by this reader. Current clue display is capped at 240 groups, with omitted counts.

Reviewer checklist

Assign an owner and status to each action and record a verification reference in Notes. Verified by reviewer is your statement, not automated proof. Do not enter credentials or sensitive values. Save the handover checklist JSON for structured records, or save readable handover notes for a colleague. These files can include private review notes. They are not uploaded, included in AI evidence requests or preserved by ordinary report exports. Clearing/closing loses unsaved annotations. Restore saved checklist accepts JSON up to 1 MiB. New checklists use an exact ZIP SHA-256 fingerprint (available for archives up to 64 MiB); repacking changes it. Review the preview and explicitly replace current notes. Legacy checklists only match package names and require independent confirmation; no exact identity is claimed. Invalid or mismatched inputs leave notes untouched.

Change impact tab

Choose a file to find code nodes assigned to that source file by the supported reader. The view explores upstream and downstream graph directions up to three hops, 120 candidate nodes per direction and 60,000 edge checks per direction. These are investigation candidates, not proven affected components. Starting nodes use explicit source-file metadata rather than cross-file evidence locations. File association does not prove runtime execution. Dynamic and omitted references can hide relationships; inferred edges can overstate them. Use source links and isolated tests before estimating or changing software.

Choosing the right tool

Reversium is a private, model-free first step for ZIP-based software handover. Specialist tools offer deeper code quality rules, security queries or architecture workflows. The public market comparison explains the documented focus of CodeBoarding, SonarQube, CodeQL and Joern, along with Reversium limitations and planned work. It is not an exhaustive market survey or product benchmark. No model or paid service is needed for the new checklist and graph views.

Compare feature maps between deliveries

Start with two reports

Analyze the current ZIP, open Compare reports and select an earlier report JSON. Both reports stay on your device. Feature-map changes appear before the file comparison. Reports without a map show comparison unavailable; missing coverage is not interpreted as removals.

Read the changes

Counts show added, removed and changed features, nodes and relationships. Stable IDs are compared. Relationship changes include endpoints, label, confidence and recorded feature membership; evidence line movement alone is ignored. A renamed identity may appear as a removal plus an addition. Analyzer improvements and differences in coverage can change maps even when software is unchanged.

Follow evidence

Each relationship card shows earlier and current source/target names, connection labels and confidence, plus a small diagram. Current evidence links open current findings. Earlier source locations are labelled and must be checked in the saved earlier report: their finding numbers must not be interpreted against the current report.

Decide what to test

Directly recorded feature membership is shown separately from a broader neighborhood, which may contain unrelated features through shared nodes. Explore these as candidates, not proven effects or failures. Inspect the source and run isolated workflow tests before changing or estimating software. No candidates found does not prove no effect.

Bounds and saving

Imports retain the 32 MiB report limit. Graphs allow up to 20,000 nodes, 40,000 edges and 1,000 features. Invalid identities, endpoints and evidence references are rejected. The view shows up to 100 changes per node/feature list and 100 relationship cards, with omitted counts. Each relationship side explores at most two hops, 120 nodes and 200 edge visits. Feature names show up to 20 per list and evidence up to four references. Save comparison notes preserves displayed results and omissions. Checklist notes are preserved, and no comparison is sent to AI.

File alerts and isolated execution

Read the traffic lights

Open File alerts or the alert-count button above the inventory. Green means no signals from the supported checks, never guaranteed safe. Yellow means possible corruption, active content, unsupported encoding/format or incomplete inspection. Red means a potentially dangerous execution pattern or filename disguise requiring review; it is not confirmed malware. Filter by color or Possible damage / invalid format, then follow evidence links. The CRC summary separates passed, failed and unverified members.

Integrity and coverage

Members within the existing 4 MiB and compression limits are read in memory through the ZIP reader to check CRC/decompression, including inventory-only members. Corrupt members are skipped with yellow alerts while other members continue. CRC detects some accidental damage, not malicious changes. Encrypted, oversized and unsupported members remain unverified. Invalid archive structures and unsafe paths are blocked. Nested member contents are not recursively checked. File type/header checks are limited; a valid CRC does not establish a valid image or document.

Suspicious signals

Selected checks flag executable headers in passive-looking filenames, document/image names with active extensions, bidirectional filename controls, download-to-interpreter pipelines, encoded PowerShell candidates and direct decode-and-execute patterns. Text-pattern checks cover selected source/script suffixes and the first 256 KiB only. Comments and examples may match. Findings retain a reason and location, not a copy of a suspicious command. This is not antivirus, a vulnerability audit, signature/reputation verification or a complete malware detector.

Can it run the ZIP safely?

The online browser edition analyzes files and never executes them. Build 0.19.0 adds an experimental native Windows QEMU pilot for one reviewed Tudela build/start profile. It first tests a blank Alpine guest, then uses a separate disposable offline guest with synthetic settings. This is a developer CLI workflow, not a general Run ZIP button. No network adapter, host shares, clipboard integration or forwarded ports are enabled. No VM guarantees zero risk; missing services and workflow acceptance remain explicit.

What to do next

Request replacement files for integrity failures. Ask the developer to explain red patterns, script capabilities and yellow gaps. Review source before installing dependencies or running scripts, and validate only with disposable data in a separate environment. Do not upload findings or private files to third-party scanning services without authorization. Reversium keeps these checks local and never automatically executes, deletes or quarantines your files.

Setup readiness and sandbox preparation

Eight practical checks

Open Setup & readiness after analysis. Review runtime declarations, dependency lockfiles, configuration templates, migrations/seeds, build/start/test clues, deployment/container clues, backup/recovery and monitoring/operations clues. Found means matching inventory names, not successful execution. Missing means no matching filename, not proof that functionality is absent. Expand inventory evidence and save the setup review plan for a developer.

Host prerequisites

The local Reversium server can check prerequisite availability from this tab. The online edition cannot inspect host programs and leaves this button disabled. Alternatively use python -m reversium --sandbox-check. This reads OS edition and executable availability only; it does not probe guest isolation, install tools or run uploaded software.

Prepare offline profiles

Run python -m reversium --sandbox-plan reports/new-sandbox-plan with a new directory name. Existing directories are refused. Output includes host-check JSON, a blank offline Windows Sandbox profile, an explicitly untested QEMU argument draft and instructions. Nothing launches automatically. No ZIP, credentials, host mappings or startup commands are included. Networking, clipboard and peripheral redirection are disabled in the Windows profile.

What is still missing

The experimental local QEMU pilot now has guest boot, selected isolation tests, supervised termination and read-only input media. A separate Tudela API tour profile now supplies disposable database/auth/storage services. General project support, a guarded browser preview and full interactive browser journeys remain missing. A passed isolation check is not proof against emulator flaws or permission to run arbitrary packages. No paid service or Windows upgrade is required; compatible hardware, existing 7-Zip and maintenance are still needed.

Troubleshooting and a practical handover

Runtime cannot start or analysis is slow

Allow the initial app/runtime download to finish and use a current desktop browser. A network interruption, unsupported browser or device memory limit can prevent analysis. Cancel and retry a smaller source-only package, or use the local application. Online analysis stops after five minutes. Do not repeatedly reload a report you have not saved.

Few findings or an empty workflow

Check that the package includes source files and manifests. Read File inventory to see which files were content-inspected versus inventoried only. Read Limitations for unsupported syntax and unresolved links. Binary-only or incomplete source deliveries cannot establish a full platform map.

An evidence link or download is unclear

A source link selects an Evidence row; it does not fetch a repository file. Downloads normally go to your browser download folder or trigger a save dialog. Check browser download controls if a file is missing. HTML is for reading; JSON is for tools; SVG is the drawn diagram; the AI request is only selected text you reviewed.

AI connection or response fails locally

Check the local assistant wizard, optional dependencies, network connection, account eligibility and provider usage limits. Review the error without posting credentials. Signing in and checking connection do not themselves share the ZIP. Do not assume a successful connection proves inference availability; there is no automatic paid API fallback. Online has no live account connection to repair.

A handover you can act on

Choose one important feature and its expected result. Save the HTML and JSON. Give the developer its workflow/evidence and a list of missing services, configuration and data. Agree on an isolated setup test and record the outcome, date and owner for every gap. Decide what is ready, what is blocked and what must be tested next. Keep the report with the exact delivery it describes.

Glossary

API: a request interface between software components. Callback: a function passed to run when an event occurs. Dependency: another package referenced by the code. Manifest: a file listing project settings or packages. Migration: a script describing a database change. Static analysis: reading files without executing the uploaded program. Evidence: the recorded source location or metadata that supports a finding. Unknown: a claim the available inspection has not established.

Experimental local QEMU sandbox

Who this is for

This developer-only Windows pilot is separate from ordinary browser analysis. Run python scripts/bootstrap_qemu_lab.py with existing 7-Zip to prepare pinned QEMU and Alpine assets. It extracts QEMU without running its installer, changing PATH or enabling Windows features. An optional pycdlib ISO library is installed only inside the private lab. Published checksums are checked over HTTPS; this is not independent signature attestation.

Provision without connecting the guest

Run python scripts/prepare_tudela_offline.py "path/to/Tudela.zip" to cache public locked libraries and Alpine packages. Package names are requested from public registries; ZIP/source contents are not uploaded. Packages are not installed or project scripts executed on the host. APK signatures and npm lockfile integrity are checked before use. Lifecycle install hooks are skipped.

Test isolation first

Run python -m reversium --sandbox-self-test. Then python -m reversium --sandbox-tudela "path/to/Tudela.zip" repeats the blank-guest checks before loading application input. TCG software emulation is the default. Add --sandbox-accelerator whpx only if Windows hypervisor support is already enabled. The same explicit accelerator is tested and used; there is no automatic networking or accelerator fallback.

Read the result honestly

The fixed Tudela profile builds the frontend and backend, starts the server with reserved .invalid service addresses and synthetic keys and requests its home page from inside the guest. This is not proof of working database, login, storage, AI or published-tour workflows. No guest browser preview is exposed on the host. This base guest has 1 GiB RAM; the separate API tour profile uses 3 GiB. Both have two virtual CPUs, bounded logs and a maximum ten-minute deadline; Windows job supervision stops it.

Private files retained

The reports/qemu-lab directory retains source-bearing read-only ISO inputs, package caches, logs and result JSON. Guests have temporary RAM-backed state and no writable persistent disk; stopping them does not erase the retained host inputs or logs. These ignored artifacts must not be committed or shared. Common environment/key files are omitted from the guest copy, but embedded secrets may still exist. Keep live credentials out of all experiments.

Link sandbox tests to the exact ZIP

Generate the right record

Use the local CLI with Reversium 0.21.0 or later: python -m reversium --sandbox-tudela "path/to/Tudela.zip". Use --sandbox-accelerator whpx only where already supported. Fresh blank-guest isolation checks run before application input. Each application attempt writes run-evidence.json in its private run directory and reports/qemu-lab/latest-run-evidence.json. Choose this portable record, not private result.json or latest-self-test.json. It contains fixed check outcomes, fingerprints and metadata, without source, commands, credentials or raw logs.

Preview, match, then attach

Analyze the original ZIP, open Setup & readiness and choose the sandbox result JSON (maximum 256 KiB). Reversium requires an exact SHA-256 match, including archive packaging. A renamed identical ZIP matches; a repacked or changed ZIP does not. Missing fingerprints, unsupported profiles and malformed records are refused. Selection previews only. Review the record, then choose Attach this reviewed test record. A rejected replacement leaves the earlier attached record unchanged.

Read what was actually tested

Recorded application outcomes and blank-guest isolation checks are shown separately from Static setup clues. Each check says recorded pass, recorded failure or not recorded. A pass describes only that bounded check. The Tudela profile tests a filtered copy, omitting common environment/key/generated files and using synthetic configuration. Its original ZIP fingerprint and tested-copy fingerprint are both displayed; this is not execution of every byte in the original delivery.

Keep assumptions and trust limits visible

The base profile leaves database, login, storage and published tours unverified. The API tour profile records selected real-service checks separately; interactive browser clicks, production operation and broader coverage remain unknown. A successful internal HTTP request does not prove the feature map or platform stability. Isolation checks do not certify a VM free of security flaws. Imported JSON and timestamps can be edited: a matching fingerprint links the claimed result to a delivery but does not authenticate the producer or establish that the claims are true. Review original local logs privately when needed.

Save, reopen or remove

Save linked test record (JSON) keeps the validated record for later attachment to the same ZIP. Save readable test summary produces local text notes. A readable report HTML export retains an explicitly attached record and validates it again when reopened; pending previews are excluded. Standard technical report JSON stays static analysis only. Remove linked test record clears the attachment. Nothing persists automatically after closing the page, changes handover checklist statuses or enters an AI request. Files and records remain on your device.

Publish a disposable Tudela tour through the API

Prepare only public service assets

After the normal QEMU/Alpine and Tudela dependency preparation, run python scripts/prepare_tudela_services.py. This caches pinned official Auth/PostgREST releases, selected Storage runtime image layers and signed Alpine PostgreSQL packages. It downloads public components; it never uploads your source or installs these services on Windows. Storage uses its real application and dependencies on the guest Node runtime, not a complete vendor Docker deployment. Checksums are publisher/image digests over HTTPS, not independent signature attestation.

Run the separate tour profile

Run python -m reversium --sandbox-tudela-tour "path/to/Tudela.zip" --sandbox-accelerator whpx where WHPX is already available. The reviewed tour profile uses a 3 GiB guest, two virtual CPUs, a bounded 2 GiB RAM workspace and a ten-minute application deadline. It repeats all isolation checks at that guest size before loading the application. Default software emulation can be slow; compatibility and local hardware still matter. No NIC, host shares, port forwarding or live service values are introduced.

What the profile actually checks

Real PostgreSQL, Auth, PostgREST and file-based Storage run inside the guest. Disposable credentials and a local TLS certificate are generated there; certificate verification remains enabled. Original Tudela migrations are applied. A reviewed derivative of the supplied acceptance script logs in separate editor/reviewer users, rejects a wrong password, checks role/tenant boundaries, uploads validated media, approves and publishes one tour with two ordered stops, then retrieves public tour JSON and the exact uploaded image bytes through Tudela HTTP routes. This is API end-to-end acceptance, not CMS button-click testing.

Read the linked result

The tour writes the same portable run-evidence.json/latest-run-evidence.json paths using a distinct v2 profile. Attach it in Setup & readiness to the exact matching ZIP. Recorded passes, failures and unrecorded checks stay separate from static findings. Prepared service versions and guest memory are shown; version metadata alone is not health evidence. Older v1 build/start records remain supported. A fingerprint match does not authenticate an editable record or certify the whole platform.

Cleanup and remaining limits

The supplied acceptance cleanup closes disposable tenants and suspends its users; guest termination discards RAM-backed database/files and stops all services. This is not secure-memory-erasure certification. Private host input ISOs, caches and logs remain in the ignored lab directory. There is no host browser preview or connection to live Tudela. Interactive CMS publishing, external email/OAuth, live AI/maps, production stability and comprehensive security remain unverified. No paid software, cloud or API service is needed; local hardware, disk and downloads are still required.

Test the CMS in a disposable browser preview

Use the separate explicit command

After preparing the QEMU lab, Tudela dependencies and service assets, run python scripts/test_tudela_cms.py "path/to/Tudela.zip" --accelerator whpx where WHPX is already available. The default is slower TCG emulation. This reviewed Tudela profile repeats the 3 GiB isolation gate and starts new services and synthetic editor/reviewer accounts. It is not an arbitrary ZIP runner or a feature of the hosted browser edition.

How the browser reaches an offline guest

The guest still has no network adapter, host shares or forwarded network port. A temporary random loopback HTTP origin relays bounded requests over the supervised serial channel to a fixed guest HTTPS gateway. Guest TLS verification remains enabled. The browser preview uses HTTP on host loopback; this is a deliberately different browser build configuration, not a production deployment. Exact Host/Origin checks, an HttpOnly SameSite entry cookie, request/response limits and a browser policy blocking external connections constrain the relay. These checks are not a security certification.

Open only the private preview entry

The stop form requires an independent CSRF confirmation token as well as the entry cookie. The command prints the path to preview-entry.json in its private run directory. Open that local entry URL on the same computer. The accompanying cms-fixture-private.json contains disposable account details for this run only; do not distribute either file or reuse real credentials. Public visitor screens and the CMS are reached through this one local origin. External maps, fonts and integrations may be blocked intentionally. Requests are serialized, so the preview can feel slower than a deployment.

Record clicks and blockers honestly

The recorded pilot passed editor wizard submission, separate reviewer approval/publication and signed-out visitor display of one tour with two ordered stops. No browser media upload, narration, maps or walking is established. Sign in as the disposable editor, create content and submit it for review. Switch to the disposable reviewer to approve and publish, then check the public visitor view. Record the screen, visible result and any validation error at each step. Do not replace a blocked UI step with a database edit and call it a browser pass. Account and tenant scaffolding is prepared before the journey; that preparation itself is not tested by clicks. Manual observations stay separate from automatic API test records.

Stop and retain only what you need

The preview allows up to twenty minutes after readiness, within a thirty-minute total guest limit, with 3 GiB guest RAM, a 2 GiB workspace, two CPUs and a bounded 32 MiB private serial log. Open /__reversium on the same preview origin and choose Stop preview and guest, or let the deadline expire. This stops the listener and guest; temporary guest database/files disappear. Source-bearing ISOs, local accounts, browser session data, screenshots and private logs can remain on the host or in the browser. Do not publish them. cms-result.json describes setup/teardown and leaves browser publishing unverified until a human records the actual observations. It is not imported as a successful API result.